

**CHIATECH JOURNAL**

## **Registered publication record**

# **Prompting for Human-Centred Output: The HUMAN Framework for Credible, Context-Aware Generative AI Writing**

Shiaondo Kenneth Chia

Volume 1, Issue 2 · Pioneer Volume · August 2026 · Part 2 · e047 · 2026-08

<https://doi.org/10.68232/cj.e047>

## **Recommended citation**

Chia, S. K. (2026). Prompting for Human-Centred Output: The HUMAN Framework for Credible, Context-Aware Generative AI Writing. CHIATECH JOURNAL, 1(2), e047.  
<https://doi.org/10.68232/cj.e047>

## **Licence and publisher**

CC BY 4.0 · Published by CHIA TECH SOLUTIONS AND RESOURCES LIMITED. ISSN assignment pending.

This bibliographic cover was added without altering the approved article pages that follow. Research content, figures, tables, references and visible source-page numbering are preserved.

# Prompting for Human-Centred Output: The HUMAN Framework for Credible, Context-Aware Generative AI Writing

**CHIA Shiaoondo Kenneth**

CHIA TECH Solutions and Resources Ltd, Abuja, Nigeria  
ORCID: 0009-0009-6434-4586 · Correspondence: chiakenneth01@gmail.com

|                      |                                                                                                                                         |
|----------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| Evidence design      | Critical review and conceptual framework; no primary dataset, participant recruitment or field experiment was conducted for this paper. |
| Planned open licence | CC BY 4.0 upon authorized publication; no licence is asserted by this production copy alone                                             |

### Abstract

Generative artificial intelligence can accelerate writing, yet a fluent draft may remain generic, unsupported or misaligned with its reader. This conceptual practice review develops the HUMAN framework for credible and context-aware AI-assisted writing: Human intent and evidence; User and audience alignment; Message architecture; Authentic detail and accountability; and Natural revision and judgement. An integrative synthesis connected research on prompt design, human-AI co-writing, iterative refinement, hallucination and educational AI governance with the practical requirements of academic, professional and public communication. The synthesis shows that quality is created both upstream, by supplying verified source material and operational constraints, and downstream, through targeted human review. Surface instructions to 'sound human' cannot substitute for evidence, a real audience, a section-level purpose and accountable verification. The paper provides a nine-part prompt architecture, an editorial rubric and adoption controls covering privacy, attribution, bias, confidentiality and disclosure. The framework does not claim to detect AI authorship or to legitimise impersonation and detector evasion. It positions generative AI as a bounded drafting technology whose output becomes publishable only through human selection, verification and approval. Experimental validation across models, genres, languages and levels of writer expertise remains necessary.

**Keywords:** prompt engineering; generative AI; human-AI collaboration; editorial quality; source grounding; responsible AI

### 1. Introduction

Generative AI has changed the speed and cost of drafting. In a controlled experiment involving common professional writing tasks, access to a generative system reduced completion time and increased average quality ratings, with larger gains for weaker initial performers (Noy & Zhang, 2023). Speed, however, is not equivalent to epistemic quality. A fluent paragraph can contain invented facts, generic examples, blurred attribution or conclusions that exceed the evidence. Writing technologies therefore need a design standard that keeps human purpose, knowledge and accountability visible.

Research on human-AI co-writing shows an interactive process: writers request suggestions, accept, reject and reshape them as their intentions evolve (Lee et al., 2022). Higher levels of task scaffolding can improve performance for some writers, but can also reduce felt ownership if the system takes too much control (Dhillon et al., 2024). The practical implication is that an effective framework must support the writer without treating generated prose as an autonomous author.

This paper develops the HUMAN framework and translates it into prompt and editorial controls. 'Human-centred' means oriented towards a real reader, grounded in supplied evidence and context, and approved

through responsible human judgement. It does not mean disguising AI use, imitating a living writer, fabricating experience or attempting to defeat AI-detection tools.

## 2. Integrative Design-Synthesis Method

The framework was developed through integrative design synthesis rather than a systematic review or experiment. Peer-reviewed research was selected on prompt patterns, non-expert prompting behaviour, co-writing, scaffolded interaction, productivity, creativity and hallucination. UNESCO guidance supplied governance principles for human agency, privacy, inclusion and age-appropriate educational use. Concepts were coded by the function they perform in writing: grounding, audience alignment, architecture, authenticity, verification and revision.

Related functions were clustered into five dimensions and translated into three practical objects: an instruction that can be included in a prompt, an observable quality criterion, and an editorial action when that criterion fails. The framework was then stress-tested conceptually against four tasks—an academic explanation, a policy brief, a teaching note and a customer-facing article—to ensure that each dimension could be expressed without relying on model-specific commands.



Figure 1. The HUMAN cycle for accountable AI-assisted writing.

Note. Original conceptual diagram developed for this paper; the sequence is iterative and has not been validated as a measurement scale.

## 3. Why Generic AI Writing Occurs

### 3.1 Under-specified tasks

A prompt such as 'write a professional article about digital education' gives a topic but not a decision, evidence base, audience, context, exclusions or acceptance test. Non-expert users often struggle to translate intentions into robust prompts and to diagnose failure after an unsatisfactory response (Zamfirescu-Pereira et al., 2023). The model then fills missing structure with statistically common patterns: broad openings, familiar lists and generic conclusions.

### 3.2 Fluency without verification

Language models generate likely continuations; they do not independently establish that a cited paper exists or that a local observation occurred. Hallucination research shows that confident false statements can arise from data, training and inference conditions (Ji et al., 2023). A prompt that requests references without supplying verified sources invites citation fabrication. A request for a personal story invites false provenance unless the writer provides a genuine account or explicitly asks for a hypothetical example.

3.3 Convergence of ideas and voice

Generative assistance can improve the judged creativity of individual outputs while reducing diversity across a collection of outputs (Doshi & Hauser, 2024). Writers therefore need to contribute distinctive questions, local evidence, experience, examples and judgements. Style variation alone is insufficient: a text may vary sentence length and still be conceptually empty.

4. The HUMAN Framework

Table 1. HUMAN dimensions and operational controls

| Dimension                               | Prompt requirement                                                                | Editorial test                                                             |
|-----------------------------------------|-----------------------------------------------------------------------------------|----------------------------------------------------------------------------|
| H — Human intent and evidence           | State the decision or communicative purpose; supply verified facts and exclusions | Every material claim traces to the source pack or is labelled as inference |
| U — User and audience alignment         | Define reader knowledge, needs, context and desired action                        | Vocabulary, examples, depth and tone fit a real reader                     |
| M — Message architecture                | Assign a job to each section and specify the reasoning sequence                   | Each section adds a distinct unit of meaning                               |
| A — Authentic detail and accountability | Use genuine examples, local constraints and attribution; prohibit fabrication     | Specificity is supportable and provenance is clear                         |
| N — Natural revision and judgement      | Review against explicit criteria; preserve meaning and evidence                   | A human verifies, edits, discloses and approves the final text             |

Note. Authors' synthesis unless otherwise stated; application requires context-specific evidence and professional judgment.

4.1 Human intent and evidence

The source pack should be separated from the instruction. It may contain results, verified policy provisions, meeting notes, product specifications, quotations with permissions or the writer's genuine examples. The prompt should identify which materials are authoritative, prohibit unsupported additions and require missing evidence to be flagged. In academic work, the model may help organise a supplied result table; it must not invent participants, measurements, significance tests or citations.

4.2 User and audience alignment

An audience definition should include knowledge, stakes, reading conditions and action. 'Write for secondary-school principals who understand school administration but not machine learning; they need to decide whether to permit AI-assisted homework after a ten-minute read' is operational. It guides explanation depth and determines whether examples should be pedagogical, legal or technical.

4.3 Message architecture

Strong prompts assign functions: open with the decision and consequence; establish verified context; explain the mechanism; compare options; state limitations; and finish with a proportionate action. Explicit architecture reduces repetition and prevents the title from being restated as an introduction. It also makes the draft easier to audit because every paragraph can be tested against its intended job.

4.4 Authentic detail and accountability

Authenticity comes from actual evidence and bounded voice, not from a list of banned phrases. A writer may supply a short sample of their own approved writing and ask the system to preserve general characteristics such as directness and formality while avoiding copied phrases. A hypothetical case must

be labelled. A quotation must be checked against the opened source. Confidential and personal data should not be entered without authority and adequate safeguards.

4.5 Natural revision and judgement

Revision should target named defects: 'remove two unsupported claims, preserve the verified figures, reduce repetition and explain the counterargument in 120 words.' Iterative refinement can improve outputs when the feedback criteria are explicit (Madaan et al., 2023). The last decision remains human: what the text means, whether the evidence is sufficient, whether disclosure is required and whether the output should be published.

5. Nine-Part Prompt Architecture

Table 2. Reusable prompt architecture

| Part              | Question to answer                                                                     |
|-------------------|----------------------------------------------------------------------------------------|
| 1. Purpose        | What decision, learning outcome or communication result must the text support?         |
| 2. Audience       | Who will read it, what do they know and what constraints shape the reading?            |
| 3. Source pack    | Which verified materials may be used, and which source prevails if they conflict?      |
| 4. Claim boundary | What must not be invented, inferred or stated without a flag?                          |
| 5. Perspective    | What defensible stance, uncertainty and counterargument must be represented?           |
| 6. Architecture   | What is the job and approximate capacity of each section?                              |
| 7. Voice          | What level of formality, directness and terminology is appropriate?                    |
| 8. Output form    | What length, headings, tables, citation style and accessibility features are required? |
| 9. Evaluation     | Which factual, ethical and editorial tests must be applied before approval?            |

Note. Authors' synthesis unless otherwise stated; application requires context-specific evidence and professional judgment.

The architecture is modular. A simple email may need only purpose, audience, facts and a review check. A journal article requires the complete set. More words do not automatically make a prompt stronger; the goal is to expose the distinctions that matter for acceptance.

6. Editorial Rubric and Governance

A ten-criterion rubric can score purpose fulfilment, audience fit, source fidelity, factual accuracy, reasoning coherence, specificity, voice consistency, concision, accessibility and ethical/provenance compliance. Each criterion may be scored from 0 (harmful or absent) to 4 (publication-ready). A total score must not compensate for a critical failure: source fidelity, factual accountability, privacy or deceptive authorship below an acceptable threshold should block release.

- Record the tool, version or access date when required by the publication policy.
- Disclose substantive use in drafting, analysis, coding or image production; never list an AI tool as an author.
- Keep confidential manuscripts, personal data and restricted information out of unapproved services.
- Test for cultural and linguistic exclusion, harmful stereotypes and one-sided framing.
- Retain the source pack, human edits and approval record for high-stakes public communication.

UNESCO's guidance for generative AI in education and research emphasises human agency, privacy, inclusion, age appropriateness and institutional validation. These principles generalise beyond schools: an organisation should adopt a tool only after evaluating the task, users, data flows, likely harms, accessibility and mechanisms for correction.

## 7. Limitations and Validation Agenda

The HUMAN framework is a conceptual synthesis, not a validated causal model. Prompt effectiveness varies by system, language, genre and user expertise. Longer instructions can conflict, and a model may ignore even a clear constraint. Human reviewers can also miss errors or over-trust fluent prose. Controlled studies should compare the full framework with shorter baselines using blinded ratings, factual-error counts, editing time, writer ownership and reader outcomes.

Research in Nigerian and other multilingual contexts should examine whether source-rich prompting improves local relevance without narrowing language diversity. Evaluation should include low-bandwidth workflows, accessibility, education levels and the cost of verification. The framework should be revised when evidence shows that a component is ineffective or burdensome.

## 8. Conclusion

Credible AI-assisted writing is not produced by asking a model to appear human. It begins with human intent and verified evidence, defines a real user, gives each section a purpose, supplies authentic detail and ends with accountable human judgement. The HUMAN framework converts those principles into a prompt architecture and editorial gate. Used transparently, generative AI can support drafting and revision; it cannot assume authorship, establish truth or approve publication.

## Declarations

### Ethics and participant involvement

Not applicable. This review/framework used publicly available documentary and scholarly sources. It involved no human participants, animals, patient material, identifiable records or intervention.

### Data availability

No primary dataset was generated or analysed. The documentary and scholarly sources used are identified in the reference list; access conditions remain those of the source publishers.

### Funding

Prepared within the publisher's journal-launch editorial research programme. No claim of grant-funded primary research is made.

### Competing interests

CHIA TECH SOLUTIONS AND RESOURCES LIMITED is the corporate author and the publisher of CHIA TECH JOURNAL. This relationship is disclosed. Publication requires independent editorial handling and an auditable peer-review and authorization record.

### Corporate author contribution and accountability

Corporate authorship is attributed to CHIA TECH SOLUTIONS AND RESOURCES LIMITED. An authorized representative must approve the final intellectual content and accept accountability before public release; this production file does not substitute for that recorded approval.

### Generative AI disclosure

Generative AI assisted search planning, source triage, drafting support, language editing and Word-production automation. It was not treated as an author or source of evidence. Every citation, claim, legal reference and editorial decision requires accountable human verification before publication.

## Figure and table provenance

The conceptual figure and synthesis tables were created specifically for this review. They contain no patient, participant, survey or third-party image data.

## References

- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? In Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (pp. 610-623). <https://doi.org/10.1145/3442188.3445922>
- Dhillon, P. S., Molaei, S., Li, J., Golub, M., Zheng, S., & Robert, L. P. (2024). Shaping human-AI collaboration: Varied scaffolding levels in co-writing with language models. In Proceedings of the CHI Conference on Human Factors in Computing Systems. <https://doi.org/10.1145/3613904.3642134>
- Doshi, A. R., & Hauser, O. P. (2024). Generative AI enhances individual creativity but reduces the collective diversity of novel content. *Science Advances*, 10(28), eadn5290. <https://doi.org/10.1126/sciadv.adn5290>
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), Article 248. <https://doi.org/10.1145/3571730>
- Lee, M., Liang, P., & Yang, Q. (2022). CoAuthor: Designing a human-AI collaborative writing dataset for exploring language model capabilities. In CHI Conference on Human Factors in Computing Systems, Article 388. <https://doi.org/10.1145/3491102.3502030>
- Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhunoy, S., Yang, Y., Welleck, S., Majumder, B. P., Gupta, S., Yazdanbakhsh, A., & Clark, P. (2023). Self-refine: Iterative refinement with self-feedback. <https://doi.org/10.48550/arXiv.2303.17651>
- Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187-192. <https://doi.org/10.1126/science.adh2586>
- UNESCO. (2023). Guidance for generative AI in education and research. <https://unesdoc.unesco.org/ark:/48223/pf0000386693>
- White, J., Fu, Q., Hays, S., Sandborn, M., Olea, C., Gilbert, H., Elnashar, A., Spencer-Smith, J., & Schmidt, D. C. (2023). A prompt pattern catalog to enhance prompt engineering with ChatGPT. <https://doi.org/10.48550/arXiv.2302.11382>
- Zamfirescu-Pereira, J. D., Wong, R. Y., Hartmann, B., & Yang, Q. (2023). Why Johnny can't prompt: How non-AI experts try (and fail) to design LLM prompts. In Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems, Article 437. <https://doi.org/10.1145/3544548.3581388>